

CS4532 - Parallel and Concurrent Programming

Take Home Lab 1

Ranasinghe W.M.D.S.S.

Wijesinghe U.M.

Step 1

Design a solution such that you can generate mMember, mInsert, and mDelete operations of each type using a given number of threads. Briefly explain your design.

Consider the following code:

```
void generate_operation_order(int *operations, double M_member, double
M_insert, double M_delete)
{
    // Calculate number of operations for each type
    int member_ops = M * M_member;
    int insert_ops = M * M_insert;
    int delete_ops = M * M_delete;

    // Fill array with exact number of each operation
    int index = 0;
    for (int i = 0; i < member_ops; i++)
        operations[index++] = 0; // 0 represents member

    for (int i = 0; i < insert_ops; i++)
        operations[index++] = 1; // 1 represents insert

    for (int i = 0; i < delete_ops; i++)
        operations[index++] = 2; // 2 represents delete

    // Shuffle the array to randomize the order of operations
    for (int i = 0; i < M; i++)
    {
        int j = rand() % M;
        int temp = operations[i];
        operations[i] = operations[j];
        operations[j] = temp;
    }
}
```

What we do here is, we denote an operation type by a given digit (0, 1, 2). Then according to the fraction, we fill an array with the required number of digits corresponding to each operation type. Finally, we can shuffle the list to randomise the order of operations.

We can then pass this array as an argument and obtain the operation type corresponding to the i^{th} operation by doing something like the following code:

```
int *operations = args->operations;
int *values = args->values;
int ops_per_thread = args->ops_per_thread;

// Perform the operations based on the pre-made operations array and
random values array
for (int i = 0; i < ops_per_thread; i++)
{
    if (operations[i] == 0)
    {
        int value = rand() % 65536; // Random value between 0 and
2^16-1
        member_rwlock(value);
    }
    else if (operations[i] == 1)
    {
        int value = values[i]; // Use the pre-generated random value
        insert_rwlock(value);
    }
    else
    {
        int value = rand() % 65536; // Random value between 0 and
2^16-1
        delete_rwlock(value);
    }
}
```

In fact, we can extend this approach for our random values too, as seen under the second if condition in the code block.

Step 2

Implement the link list using the above three approaches.

Implementation available in the code (zip file).

Step 3

Run your program a sufficient number of times under the following three cases, and then fill up the given tables using the execution time of your programs. While determining the execution time, consider only the time taken to perform m operations and ignore other overheads such as initially populating the link list.

Specs of the machine used

Device name LAPTOP-4AOFPDB3

Processor Intel(R) Core(TM) i7-10510U CPU @ 1.80GHz 2.30 GHz

Installed RAM 16.0 GB (15.9 GB usable)

System type 64-bit operating system, x64-based processor

Pen and touch No pen or touch input is available for this display

Edition Windows 11 Home Single Language

Version 23H2

Installed on 4/6/2023

OS build 22631.4169

Experience Windows Feature Experience Pack 1000.22700.1034.0

| | | | |
|-------------|----------|---------------------|----------|
| Utilization | Speed | Base speed: | 2.30 GHz |
| 20% | 2.75 GHz | Sockets: | 1 |
| Processes | Threads | Handles | Cores: |
| 340 | 6165 | 266639 | 4 |
| | | Logical processors: | 8 |
| Up time | | Virtualization: | Enabled |
| 23:17:34:24 | | L1 cache: | 256 KB |
| | | L2 cache: | 1.0 MB |
| | | L3 cache: | 8.0 MB |

| | | | |
|---------------------|----------------|--------------------|-----------|
| In use (Compressed) | Available | Speed: | 2667 MT/s |
| 13.5 GB (369 MB) | 2.3 GB | Slots used: | 2 of 2 |
| Committed | Cached | Form factor: | SODIMM |
| 23.1/30.4 GB | 2.4 GB | Hardware reserved: | 149 MB |
| Paged pool | Non-paged pool | | |
| 1000 MB | 979 MB | | |

Utilization

11%

Shared GPU memory

0.7/7.9 GB

Driver version:

27.20.100.9268

Driver date:

2/5/2021

DirectX version:

12 (FL 12.1)

Physical location:

PCI bus 0, device 2, function 0

GPU Memory

0.7/7.9 GB

Utilization

0%

Dedicated GPU memory

0.0/2.0 GB

Driver version:

30.0.15.1169

Driver date:

2/1/2022

DirectX version:

12 (FL 12.1)

Physical location:

PCI bus 2, device 0, function 0

Hardware reserved memory:

64.3 MB

GPU Memory

0.0/9.9 GB

Shared GPU memory

0.0/7.9 GB

GPU Temperature

57 °C

60 seconds

Active time

2%

Average response time

0.2 ms

Capacity:

954 GB

Formatted:

929 GB

System disk:

Yes

Page file:

Yes

Type:

SSD

Read speed

371 KB/s

Write speed

156 KB/s

Determining sample size

Consider the following formula we obtained by running a maths source through chatGPT, as well as the pre-calculated table of values we found on a separate maths source. Based on these, we found the number **385** to be required for a confidence interval of 95% with an accuracy of $\pm 5\%$.

There is an interesting trade-off between the level of confidence and the sample size that shows up here when considering the cost of sampling. The table below shows the appropriate sample size at different levels of confidence and different margins of error, assuming $p = 0.5$. Looking at each row, we can see that for the same margin of error, a higher level of confidence requires a larger sample size. Similarly, looking at each column, we can see that for the same confidence level, a smaller margin of error requires a larger sample size.

| Required Sample Size (90%) | Required Sample Size (95%) | Margin of Error |
|----------------------------|----------------------------|-----------------|
| 1691 | 2401 | 2% |
| 752 | 1067 | 3% |
| 271 | 384 | 5% |
| 68 | 96 | 10% |



To calculate the number of samples you need for a certain margin of error ($\pm 5\%$) and confidence level (95%), you can use the following formula for sample size n :

$$n = \left(\frac{Z^2 \cdot p \cdot (1 - p)}{E^2} \right)$$

Where:

- Z is the Z-score, which corresponds to the confidence level (for 95% confidence level, $Z \approx 1.96$),
- p is the estimated proportion of the population (if unknown, you can use 0.5 for the worst-case scenario, as it maximizes sample size),
- E is the margin of error (in this case, 5%, or 0.05).

Let's calculate the sample size for 95% confidence level and $\pm 5\%$ margin of error, assuming $p = 0.5$:

$$n = \left(\frac{(1.96)^2 \cdot 0.5 \cdot (1 - 0.5)}{(0.05)^2} \right)$$

Let me compute that for you.

The required sample size for a $\pm 5\%$ margin of error with a 95% confidence level is approximately 384 samples. [5.2](#)

Performance Tables

Case 1

$n = 1000$, $m = 10000$, $mMember = 0.99$, $mIndert = 0.005$, $mDelete = 0.005$

| Implementa tion | No of Threads | | | | | | | |
|--------------------|---------------|----------|----------|----------|----------|----------|----------|----------|
| | 1 | | 2 | | 4 | | 8 | |
| | Avg | Std | Avg | Std | Avg | Std | Avg | Std |
| Serial | 0.041958 | 0.006542 | | | | | | |
| One mutex | 0.041696 | 0.007086 | 0.063327 | 0.00942 | 0.069481 | 0.01418 | 0.071169 | 0.009621 |
| Read Write lock | 0.041652 | 0.004326 | 0.025719 | 0.004007 | 0.018852 | 0.007239 | 0.016675 | 0.00359 |

Case 2

$n = 1000$, $m = 10000$, $mMember = 0.90$, $mIndert = 0.05$, $mDelete = 0.05$

| Implementa tion | No of Threads | | | | | | | |
|--------------------|---------------|----------|----------|---------|----------|----------|----------|----------|
| | 1 | | 2 | | 4 | | 8 | |
| | Avg | Std | Avg | Std | Avg | Std | Avg | Std |
| Serial | 0.053418 | 0.007078 | | | | | | |
| One mutex | 0.053034 | 0.006115 | 0.077901 | 0.01262 | 0.079457 | 0.009774 | 0.081569 | 0.012634 |
| Read Write lock | 0.054478 | 0.007456 | 0.05113 | 0.00864 | 0.044636 | 0.006535 | 0.050457 | 0.010017 |

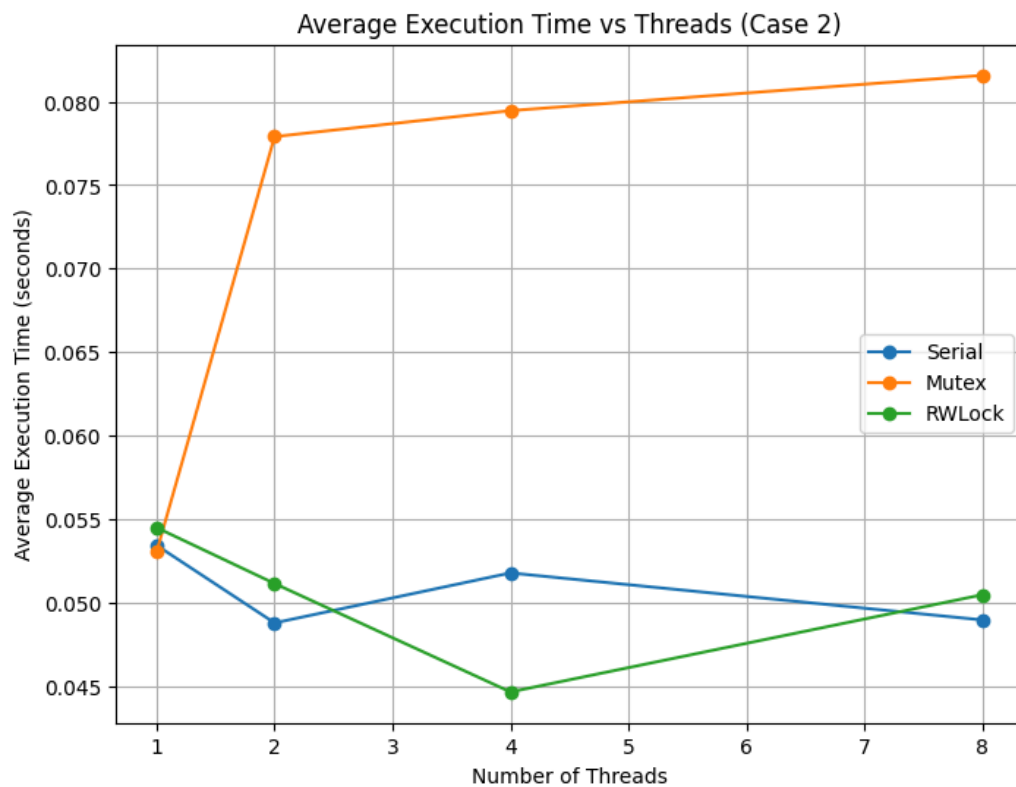
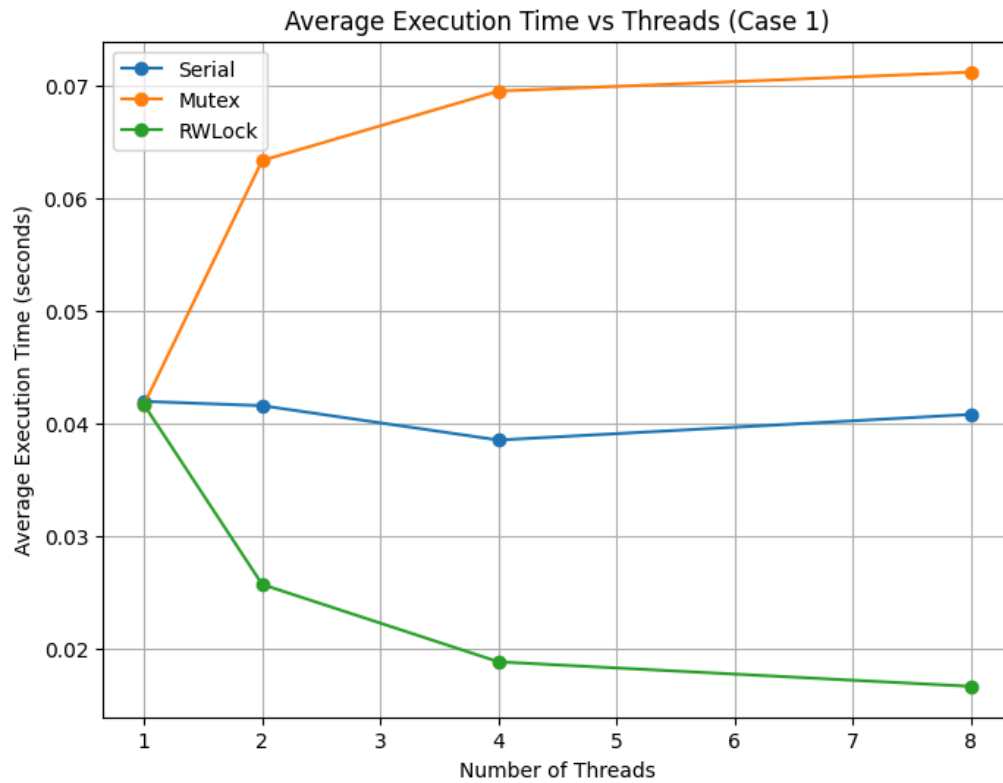
Case 3

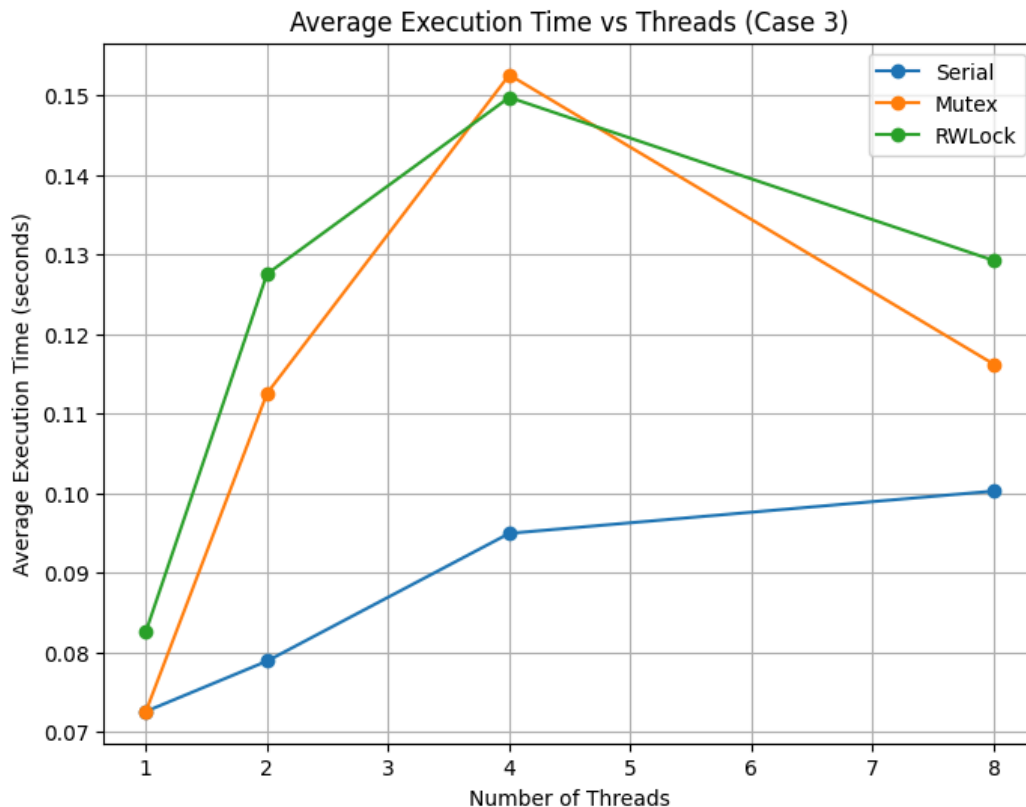
$n = 1000$, $m = 10000$, $mMember = 0.50$, $mIndert = 0.25$, $mDelete = 0.25$

| Implementa tion | No of Threads | | | | | | | |
|--------------------|---------------|----------|----------|----------|----------|----------|----------|----------|
| | 1 | | 2 | | 4 | | 8 | |
| | Avg | Std | Avg | Std | Avg | Std | Avg | Std |
| Serial | 0.072595 | 0.010607 | | | | | | |
| One mutex | 0.072603 | 0.008715 | 0.112556 | 0.0248 | 0.152564 | 0.048895 | 0.116221 | 0.021932 |
| Read Write lock | 0.082631 | 0.01573 | 0.127556 | 0.023525 | 0.14974 | 0.027703 | 0.129249 | 0.027284 |

Step 4

Plot the average execution time against the number of threads for each case.





Step 5

Comment your observations while critically evaluating your findings. You may plot additional graphs to support your discussion.

Based on the initial graphs we can see that the average execution time for each method when the number of threads are 1 is almost the same, except for case 3, where the overhead associated with the read-write lock for write operations slightly elevates it. Taking a look at the times taken for the other two methods in comparison to the serial method, we can make the following observations.

In case one, where the overwhelming majority of the operations are read operations, we can see that the time taken by the mutex method increases considerably before flattening with the number of threads. This is explained since there is no real gain for the mutex method since no two threads can actually work at the same time, taking at least as much time as the serial method just for the operations, which is further exacerbated by the overhead coming both with thread creation and possibly context switching as well as overheads for mutex acquisition times, and possibly even lower level factors such as memory contention and L1 caching in the cpu cores.

In contrast, we see the rw-lock method significantly decreasing in runtime with the number of threads before flattening out, which is explained by the parallelism which happens for read operations, since multiple threads can acquire the lock to read at the same time. While

theoretically the performance gain should increase with the number of threads consistently, the flattening out is explained by the practical overheads of thread creation, thread context switching as well as lock acquisition times and again possibly even lower level factors such as memory contention and L1 caching in the cpu cores.

In cases two and three, we can see that the performance advantage that the rw-lock method had has reduced as the fraction of write operations increases. This is explained by the fact that only the read operations can be done in parallel by multiple threads at the same time, whereas the write lock can only be held by a single thread. In fact, we see that for most thread values, rw-lock method takes slightly more time than the mutex method for case 3, indicating a slightly higher overhead for write lock acquisitions under the rw-lock method. It may be necessary to verify the idea of a higher overhead by increasing the fraction of write operations even more (maybe to 0.9) (We haven't done this in this assignment).

In addition we can attribute the extra variation within the plots to system resource availability and other system related factors. For example, there is a slight and sometimes considerable variance for the time taken to run the serial operations which shouldn't exist theoretically. It may be necessary to reaffirm this conjecture by creating a separate program which only runs one method at a time (serial, mutex or rw-lock) (Again, we haven't done this in this assignment).

